

# Notes On TensorFlow & Keras

Ara Patapoutian, 2017

## CONTENTS

<b>I</b>	<b>Introduction</b>	1
I-A	Graphs . . . . .	2
I-B	Environments . . . . .	2
<b>II</b>	<b>Tensors</b>	4
<b>III</b>	<b>Operators</b>	5
<b>IV</b>	<b>Session &amp; Train</b>	8
IV-A	Session . . . . .	8
IV-B	Train . . . . .	9
<b>V</b>	<b>Image Processing</b>	11
V-A	MNIST . . . . .	11
<b>VI</b>	<b>Keras</b>	11

## I. INTRODUCTION

- General:
  - *TensorFlow*, or TF, is an open source software library for numerical computation, machine learning & NN
  - TF is an *application program interface* (API) which is a set of routines, protocols, and tools for building software applications
    - \* the core of TF is written in C++
  - the primary API is accessed through Python
    - \* the variable that displays version `tf.__version__`
    - \* TF can also be accessed through the C & C++
  - TF has better support for distributed systems than competing software packages
- TF includes three packages
  - TensorFlow
  - TensorBoard
  - TensorFlow Serving
- TensorBoard:
  - *TensorBoard* is a visualizations software that is included with TF installation
- TensorFlow Serving:
  - *TF Serving* is a software that facilitates easy deployment of pre-trained TF models
  - a user can export their model to a file which can then be read by TF serving
  - TF Serving seamlessly switches old models with new ones
  - TF Serving enables users to avoid re-implement their models for production
  - TF Serving is is written in C++

## A. Graphs

- Computational graph:
  - TF uses *data-flow graphs* or *computational graph*
  - a computational graph is a directed graph
  - no computation is performed while a graph is constructed
  - graph representation allows distributing computations across multiple CPU/GPUs
  - a NN is naturally modeled on a computational graph
- Nodes and edges:
  - nodes constitute operations or *ops*
    - \* the nodes are objects of type `tf.Operation`
    - \* each node is differentiable
  - edges represent data by tensors
    - \* the edges are objects of type `tf.Tensor`
    - \* only tensors may be passed between nodes
    - \* hence the name TensorFlow
- Differentiation:
  - differentiation is used for gradient based optimization or training
  - TF can take the derivative of any node with respect to any other
  - as a rule of thumb, back propagation takes about twice the memory and twice the computation as the forward path
- Graph
  - graphs belong to the class `Graph`
  - `default_graph = tf.get_default_graph()`
    - \* `get_default_graph()` provides a handle of the default graph
    - \* a default `Graph` is created when TF is loaded
  - `g = tf.Graph()` constructs a graph explicitly
  - `g.as_default`
    - \* to add an operation to graph `g`, use `g.as_default` with `g.as_default()`:
 

```
a = tf.mult(2,3)
```
    - \* see below for more on `with` statement
  - `w1 = graph.get_tensor_by_name("w1:0")`
    - \* returns the `Tensor` with the given name
    - \* allows access to saved variables, tensors, or placeholders
  - `tf.reset_default_graph()` resets a graph
- with
  - the `with` statement is used to wrap the execution of a block with methods defined by a context manager
  - a *context manager* is an object that defines the runtime context to be established when executing a `with` statement
  - the context manager handles the entry into, and the exit from, the desired runtime context for the execution of the block of code

## B. Environments

- TF through `anaconda`
  - go to the desired directory
  - to activate
 

```
source activate env_name
```

    - \* my environment names are `tensorflow` and `carnd-term1`
  - choose `python` or `jupyter notebook`
  - to deactivate
 

```
source deactivate
```
- TF through `virtualenv`
  - to activate type
 

```
source ~tf_env/bin/activate
```
- TF through Python:

- in virtualenv, enter Python by typing  
`python`
- from the Python cursor type  
`import tensorflow as tf`
- to avoid overhead, TF describes a graph of interacting operations that run entirely outside Python
- the role of the Python code is to build this external computation graph, and to dictate which parts of the computation graph should be run

## II. TENSORS

- Variables in TF & Python:
  - names in Python-space are temporary pointers to TF variables during the run of the script
  - saving & restoring variables are done through TF namespace, not Python
- Tensor:
  - the *tensor* data-structure is used to represent all data
  - a tensor has a dynamic rank, shape & a static type
    - \* *rank* is the number of dimensions of the tensor
    - \* *shape* describes the the sizes of each dimension
  - list of static TF data *types* and the corresponding Python *dtype*:

TF	Python	bits
DT_FLOAT16	tf.float16	16
DT_FLOAT	tf.float	32
DT_DOUBLE	tf.double	64
DT_INT8	tf.int8	8
DT_INT16	tf.int16	16
DT_UINT8	tf.uint8	8
DT_UINT16	tf.uint16	16
DT_QINT8	tf.qint8	8
DT_STRING	tf.string	8*
DT_BOOL	tf.bool	1

\* letters U & Q respectively indicate unsigned and quantized

- `tf.Variable()`
  - a `Variable()` is a value that lives in TensorFlow's computation graph
    - \* the value returned by `tf.Variable()` is an instance of the Python class `tf.Variable`
  - a TF variable reside in memory buffer, containing tensors
    - \* variables can also be saved to disk
    - \* variables are mutable tensors
  - the value of a `Variable()` can be changed using the `Variable.assign()` method
    - \* they are executed in `Session`, see Section IV

```
b = a.assign(a*2)
```
  - addition & subtraction can be performed using specific methods
 

```
b = a.assign_add(2)      #b=a+2
c = a.assign_sub(1)      #c=a-1
```
  - if a variable should not change during training, define it as
 

```
a = tf.Variable(1, trainable=False)
```
  - calling `tf.Variable()` adds several operations to the graph:
    - 1) a variable operation that holds the variable value
    - 2) an initializer operation that initializes the variable
      - \* initialization is a `tf.assign()` operation
    - 3) the ops for the initial value (e.g. zeros op), are also added to the graph
- Special tensors:
 

```
tf.zeros(shape, dtype=tf.float32)
    e.g., tf.zeros([3, 4], tf.int32)
tf.zeros_like(tensor, dtype=None)
tf.ones(shape, dtype=tf.float32)
tf.ones_like(tensor, dtype=None)
tf.fill([3,4], value, name=None)
tf.constant(value, dtype=None, shape=None)
    e.g., tf.constant([4, 2, 3])
tf.linspace(start, stop, num)
tf.range(start, limit=None, delta=1)
tf.random_normal(shape, mean=0.0, stddev=1.0, dtype=tf.float32, seed=None)
tf.truncated_normal(shape, mean=0., stddev=1.0, dtype=tf.float32, seed=None)
```

```
tf.random_uniform(shape, minval=0, maxval=None, dtype=tf.float32, seed=None)
tf.random_shuffle(value, seed=None) #along first dimension
tf.random_crop(tensor, size, seed=None) #randomly crops a tensor to a given size
```

- Random seeding:
  - random operations use two seeds: graph-level & operation-level seeds
  - the following command sets the graph-level seed
 

```
tf.set_random_seed(seed)
```
- `tf.Variable()` Initialization:
  - variables proceed through 3 steps:
    - \* specified
    - \* added, or initialized, to graph
    - \* executed
  - when a variable is created, a tensor is specified as its initial value to the `Variable()` constructor
  - example:
 

```
* x = tf.Variable([1.0, 2.0], name="y")
* tf.Variable() only specifies the initial value, it does not perform the initialization
* the state of a variable is managed by class Session, see Section IV
* a Variable should be initialized within a Session
```
  - special tensors can be used to initialize a `tf.Variable()`

```
a = tf.Variable(tf.zeros([3, 4], tf.int32))
```
  - to initialize a variable from the value of another variable, use `initialized_value()` method

```
w2 = tf.Variable(w1.initialized_value())
```
  - an op that can be used within a `Session`, to initialize all the variables is

```
init_op = tf.global_variables_initializer()
```

    - \* note that the above command adds the op the graph, it still does not run it
    - \* to initialize need to use `run()` command within session
  - an op that is used to initialize some of the variables is

```
init_op = tf.initialize_variables([v1])
```
  - variable initializers must be run explicitly, before other operations can be run
- `tf.device()`
  - a variable can be pinned to a particular device,
  - examples of devices are CPUs & GPUs,
 

```
tf.device("/cpu:0"):
    v = tf.Variable(...)
```
- `tf.placeholder()`
  - `tf.placeholder()` is a variable that will be assigned data, at a later date
  - `tf.placeholder()` creates an input node on graph
  - examples:
 

```
a = tf.placeholder(tf.float32, name="my_input")
b = tf.placeholder(tf.float32, shape=[None, 784])
```
  - the Python data type argument (`dtype`) is required
  - `shape` is an optional parameter

### III. OPERATORS

- `tensor.get_shape()`
  - `print(x.get_shape())`
  - `x.get_shape()` is the static shape of `x`
  - the dynamic shape (during a session) can be accessed using `x.shape`
- `tf.TensorShape.as_list()`
  - returns a list of integers or `None` for each dimension
  - `x.get_shape().as_list()`

- `tf.reshape()`
  - `reshape(tensor, shape)`
  - `tf.reshape(x, [-1, 28, 28, 1])`
- `tf.transpose()`
  - `transpose(x, perm=None)`
  - transposes `x` and permutes the dimensions according to `perm`
- `tf.one_hot()`
  - `tf.one_hot(x, _size)`
  - converts `x` to one-hot format
- `tf.split()`
  - `split(value, num_or_size_splits, axis=0)`
  - splits a tensor into sub tensors
  - `num_or_size_splits` can be an integer that evenly divides `value.shape[axis]`
  - `num_or_size_splits` can also be a tensor that partitions `value` into `len(size_splits)` pieces
  - `split0, split1, split2 = tf.split(value, [4, 15, 11], 1)`
- `tf.pack()`
  - `tf.pack([a, b, c])`
  - packs a list of rank-`R` tensors into 1 rank-`(R+1)` tensor
- `tf.assign()`
  - `tf.assign(x_ref, _value)`
  - update `x_ref` by assigning `_value` to it
- `tf.cast()`
  - `tf.cast(x, tf.float32)`
- General operator examples:
 

<code>tf.add(x, y)</code>	#addition
<code>tf.sub(x, y)</code>	#subtraction
<code>tf.div(x, y)</code>	#elementwise integer or floating point division
<code>tf.truediv(x, y)</code>	#elementwise floating point division
<code>tf.mul(x, y)</code>	#multiplication
<code>tf.matmul(x, W)</code>	#matrix multiplication
<code>tf.mod(x, y)</code>	# <code>x % y</code>
<code>tf.pow(x, y)</code>	# <code>x.^y</code>
<code>tf.log(x)</code>	#natural log
<code>tf.reduce_mean()</code>	#take the average
<code>tf.reduce_sum()</code>	#take the sum
<code>tf.argmax(X, axis)</code>	#index of the highest entry along some axis
<code>tf.squared_difference(x, y)</code>	#elementwise
<code>tf.sigmoid(x)</code>	#elementwise sigmoid fn
<code>tf.equal(x, y)</code>	#1 if equal
<code>tf.less(x, y)</code>	#<
<code>tf.less_equal(x, y)</code>	#<=
<code>tf.greater(x, y)</code>	#>
<code>tf.greater_equal(x, y)</code>	#>=
<code>tf.logical_and(x, y)</code>	#&
<code>tf.logical_or(x, y)</code>	#
<code>tf.logical_xor(x, y)</code>	#^
- NN specific operators:
  - `tf.nn.max_pool(x, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')`
  - `tf.nn.relu(y)`
  - `tf.nn.tanh(y)`
  - `tf.nn.softmax(y)`
  - `tf.nn.dropout(h_fc, keep_prob)` #it multiplies kept units by `1/keep_prob`.
  - `tf.nn.sigmoid_cross_entropy_with_logits(y, y_)` #cross entropy with sigmoid
  - `tf.nn.softmax_cross_entropy_with_logits(y, y_)` #cross entropy with softmax
  - `tf.nn.sparse_softmax_cross_entropy_with_logits(y, y_)` #sparse/one-shot

```

tf.nn.bias_add(n_layer,bias)| #adds bias, addition
y = tf.nn.xw_plus_b(_x, _w, _b) # X W + B
_v, _i=tf.nn.top_k(input, k=3) #returns k highest values and their indices

```

- Image input format:

- TF's input pipeline format is optimized to work with a batch of images
- a rank-four tensor is used as the NN input structure with shape  
[in\_batch\_size, in\_height, in\_width, in\_channels]
- 1) in\_batch\_size is the batch size
- 2) in\_height is the rows of an image
- 3) in\_width is the columns of an image
- 4) in\_channels is the color channels

- tf.layers

- a library that provides a set of high-level neural networks layers
- e.g. Conv2d, Conv2d\_transpose, Dense, MaxPooling2D, etc.

- out=tf.nn.conv2d(input\_batch, filter, strides, padding)

- input\_batch is rank-4 tensor as described above
  - \* data type is float16, float32 or float64
- filter is the kernel which is a tensor with the same type as input\_batch
  - \* shape is [filter\_height, filter\_width, in\_channels, out\_channels]
- strides causes the filter to skip over pixels
  - \* a list of integers, of length 4, corresponding to the four dimensions of input\_batch
  - \* e.g strides = [1 , 2, 2, 1]
- padding, which is assigned to 'SAME' or 'VALID', describes how to fill the missing area in image
  - \* if padding='SAME', then

```

out_height = ceil(float(in_height) / float(strides[1]))
out_width  = ceil(float(in_width) / float(strides[2]))
pad_along_height = ((out_height - 1) * strides[1] +
                    filter_height - in_height)
pad_along_width = ((out_width - 1) * strides[2] +
                    filter_width - in_width)
pad_top = pad_along_height / 2
pad_left = pad_along_width / 2

```

- \* if padding='VALID', then

```

out_height = ceil(float(in_height) / float(strides[1]))
out_width  = ceil(float(in_width) / float(strides[2]))
pad_along_height = ((out_height - 1) * strides[1] +
                    filter_height - in_height)
pad_along_width = ((out_width - 1) * strides[2] +
                    filter_width - in_width)
pad_top = pad_along_height / 2
pad_left = pad_along_width / 2

```

- in detail, with the default NHWC format, the function computes

```

output[b, i, j, k] =
    sum_{di, dj, q} input[b, strides[1] * i + di, strides[2] * j + dj, q] *
    filter[di, dj, q, k]

```

- conv2d\_transpose

- sometimes called deconvolution after Deconvolutional Networks, but is actually the transpose (gradient) of conv2d rather than an actual deconvolution
- a stride of 2, and "SAME" padding would double the output dimensions

- tf.metrics

- metric functions in this library return a Tensor for the metric result and a Tensor Operation to generate the result
  - \* need to run operation before getting the result
- the other characteristic of TensorFlow metric functions is the usage of local TensorFlow variables

- \* these are temporary TensorFlow Variables that must be initialized by running `tf.local_variables_initializer()`
- examples include
  - `accuracy()`
  - `mean_squared_error()`
  - `root_mean_squared_error()`
  - `mean_iou()`

## IV. SESSION & TRAIN

### A. Session

- `tf.Session()` & `tf.InteractiveSession()`
  - sessions are responsible for graph execution
  - Tensorflow variables are only alive inside a session
  - launch one of two classes `tf.Session()` or `tf.InteractiveSession()`
  - `tf.InteractiveSession()` class makes TensorFlow more flexible about how it is structured
    - \* it allows to interleave operations that build a computation graph with ones that run the graph
  - these classes make the connection to the highly efficient C++ back-end to do its computation
  - will focus on `tf.Session()`
- `sess = tf.Session()`
  - `tf.Session()` is a constructor for a session
  - `tf.Session()` takes three *optional* arguments:
    - \* `target` specifies the execution engine
    - \* `graph` specifies the Graph object
    - \* `config` allows users to specify options to configure the session
  - the constructor places the graph operations onto specific devices
  - the associated graph attribute is `sess.graph`
- `sess.run()`
  - two common arguments of the `run()` method are
    - \* `fetches`
    - \* `feed_dict`
  - `fetches` is not optional
- `sess.run(fetches)`
  - `fetches` can be any graph element
    - \* in other words, `fetches` can be an operation or a tensor
  - if `fetches` is a tensor, the output will be a NumPy array
    - `output = sess.run(v)`
    - \* computes tensor `v` and assigns it to `output`
    - \* it is also possible to pass a list of graph elements
      - `sess.run([v, u])`
      - \* if the argument is a list so will be the output
  - if `fetches` is an operator the output will be `None`
    - \* `sess.run(tf.global_variables_initializer())` initializes all the variables
- `feed_dict`
  - `feed_dict` is an optional argument to `run()`
    - \* it overrides tensor values in the graph
  - `feed_dict` expects a Python dictionary object as input
    - \* the keys in the dictionary are handles to tensor objects that should be overridden
    - \* the values can be numbers, strings, lists or NumPy arrays
  - `sess.run(optimizer, feed_dict={x: batch_x, y: batch_y})`
    - \* in the above example, the session runs an element called `optimizer` for `x = batch_xs` & `y = batch_ys`
- `tensor.eval()`
  - `t.eval()` is a shortcut for calling `tf.get_default_session().run(t)`



- can also use it to print tensors  
`print (t.eval())`
- `sess.close()`
  - release a Session by typing `sess.close()`
- `with`
  - there are two ways of using `with` to run a Session:
  - the first approach is as follows
 

```
sess = tf.Session()
with sess.as_default():
    a.eval()
    ...
sess.close()
```
  - the second approach is
 

```
with tf.Session() as sess:
    ...
```

in this second approach, there is no need to `close()` a Session at the end

## B. Train

- `tf.train`
  - `tf.train.Optimizer` is a base class for gradient based optimizations
  - `tf.train.SummaryWriter` is a class that creates an event file in a given directory and adds summaries to it
  - `tf.train.Saver` saves the environment in proprietary format
  - `tf.train.string_input_producer()` to read files
- `tf.train.Optimizer`
  - TF uses automatic differentiation to find the gradients of the loss function with respect to each of the variables
  - specific classes include
 

```
GradientDescentOptimizer
AdamOptimizer
AdagradOptimizer
MomentumOptimizer
```
  - examples:
 

```
optimizer = tf.train.GradientDescentOptimizer(0.5).minimize(total_loss)
optimizer = tf.train.AdamOptimizer(1e-4).minimize(total_loss)
```
  - in the above two examples, new operations are added to the graph
  - the optimizer returns a single operation `optimizer`
  - when run, this operation performs a step of training
 

```
training_operation = optimizer.minimize(cost_function)
sess.run(training_operation, feed_dict={x: batch[0], y_: batch[1]})
```
- `tf.stop_gradient()`
  - `tf.stop_gradient()` prevents the gradient from flowing backwards past a specified point
  - `layer_n = tf.stop_gradient(layer_n)`
- `tf.train.SummaryWriter`
  - used by TensorBoard
  - `writer = tf.train.SummaryWriter('./my_graph', sess.graph)`
  - graph description will be stored in directory `'./my_graph'`
  - `SummaryWriter` will output the description of the graph, i.e. `sess.graph`, into `my_graph` directory
- `tf.train.Saver()`
  - the `Saver` class adds operations to save and restore variables to and from checkpoints
    - \* *checkpoints* are binary files in a proprietary format which map variable names to tensor values.
    - \* the role of the class `Saver` is to save the tf environment
    - \* `Saver` class provides the functionality to save any `tf.Variable` to a file system

- the constructor is
 

```
saver = tf.train.Saver() or
saver = tf.train.Saver({'v1': v1, 'v2': v2}) where specific variables are passed as dictionary, or
saver = tf.train.Saver([v1, v2]), where variables are passed as a list
```
- `saver.save(sess, "mymodel", global_step=3)`
  - \* the save method should be run within a Session
  - \* multiple filenames will be created or updated under the name `mymodel-3`
  - \* by default, the saver will keep only the most recent five files
- four files are generated with save method:
  - \* `mymodel-3.index` is a checkpoint file
  - \* `mymodel-3.data` is a checkpoint file that contains the values of training variables
  - \* checkpoint file has the list of checkpoint filenames
  - \* `mymodel-3.meta` is a meta graph protocol buffer which saves the complete Tensorflow graph
- to save a subset of the variables, specify the names and variables to be saved by passing a Python dictionary: keys are the names to use, values are the variables to manage,
 

```
v2 = tf.Variable(..., name="v2")
saver = tf.train.Saver({"my_v2": v2})
```
- to verify if a checkpoint is available, use command
 

```
ckpt = tf.train.get_checkpoint_state(os.path.dirname(__file__))
```
- to recreate the model (not values), type
 

```
saver = tf.train.import_meta_graph('mymodel-3.meta')
```
- to restore the variable values, type
 

```
saver.restore(sess, "/tmp/model.ckpt"), or
saver.restore(sess, tf.train.latest_checkpoint('.'))
```

  - \* while restoring Session should be active
  - \* all restored variables should be pre-defined in Python
- `saved_model`
  - to save or restore a whole model it is recommended using `saved_model`
  - `saved_model` is a language-neutral, that enables higher-level systems and tools to produce, consume, and transform TensorFlow models
  - `saved_model.loader` can load & restore pre-existing models
  - `saved_model.loader.load(sess, [tag_constants.TRAINING], export_dir)` where
    - \* `sess` is the session in which the graph definition & variables are restored to
    - \* `tag_constants.TRAINING` are the tags used to identify the MetaGraphDef to load
    - \* `export_dir` is the location (directory) of the SavedModel
- Testing it two ways:
  - `tensor_name.eval(feed_dict={x: mnist.test.images, y_: mnist.test.labels})`
  - `sess.run(tensor_name, feed_dict={x: mnist.test.images, y_: mnist.test.labels})`
- Reading data:
  - there are three main methods of getting data into a TensorFlow program
    - 1) through `feed_dict` as described in `Session()`
    - 2) reading from files, where an input pipeline reads the data from files at the beginning of a TF graph
    - 3) for small data sets can use preloaded data, where a constant or variable in the TF graph holds all the data
- Reading from files:
  - `tf.train.string_input_producer()`
    - \* pass the list of filenames to the `tf.train.string_input_producer()` function
    - \* `string_input_producer` creates a FIFO queue for holding the filenames until the reader needs them
  - `filename_queue = tf.train.string_input_producer(["file0.csv", "file1.csv"])`
  - readers:
    - \* readers convert string work-units into records
    - \* *work-units* are typically filenames, while *records* are (key, value) pairs extracted from the contents of those files
      - the key would be the associated filename
    - \* to read text files use `TextLineReader()` operation
 

```
reader = tf.TextLineReader() #constructor
```

- key, value = reader.read(filename\_queue)
- \* to convert comma-separated value (CSV) format to tensors, use `tf.decode_csv()` operation
  - `tf.decode_csv(value, record_defaults)`
    - `record_defaults` is a list of Tensor objects with types from `float32`, `int32`, `int64`, `string`
    - returns a list of Tensor objects with same type as `record_defaults`
- \* `tf.train.shuffle_batch()` creates a batch by randomly shuffling tensors

## V. IMAGE PROCESSING

- `tf.image.resize_images()`
  - `resize_images(images, size)`
  - `resized = tf.image.resize_images(x, (227, 227))`

### A. MNIST

- MNIST repository in TF:
  - MNIST data sits in the `~/MNIST_data/` directory,
    - \* `'train-images-idx3-ubyte.gz'` train
    - \* `'train-labels-idx1-ubyte.gz'` train
    - \* `'t10k-images-idx3-ubyte.gz'` test
    - \* `'t10k-labels-idx1-ubyte.gz'` test
- `input_data.py`
  - `input_data.py` imports `read_data_sets()` function from `mnist.py` file
  - \* `input_data.py` also includes multiple import commands
  - `input_data.py` is in directory `~/tensorflow/tensorflow/examples/tutorials/mnist`
  - first load the program
    - `from tensorflow.examples.tutorials.mnist import input_data`
  - `mnist.py` is in directory `~/tensorflow/tensorflow/contrib/learn/python/learn/datasets`
- `read_data_sets()`
  - to execute `read_data_sets()`, type
    - `mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)`
  - the complete command is as follows
    - `read_data_sets(train_dir, fake_data=False, one_hot=False, dtype=dtypes.float32, reshape=True, validation_size=5000)`
- Datasets & Dataset
  - `mnist` is a user defined `collections.namedtuple()` called `Datasets`, that contains three instantiations of a class called `Dataset`,
    - \* `Datasets.train`
    - \* `Datasets.test`
    - \* `Datasets.validate`
  - `Dataset`
    - \* the `Dataset` class contains two NumPy arrays: `images` and `labels` both of type `numpy.ndarray`,
    - \* each image is of size 784,
    - \* each label is of size 10.
    - \* `Dataset` class includes a method that can iterate through mini-batches,
      - `batch_xs, batch_ys = mnist.train.next_batch(batch_size)`

## VI. KERAS

- *Keras* is a higher level wrapper around Tensorflow and other tools
- Models:
  - there are two types of models available in Keras:
    - \* the *Sequential model* class and
    - \* the *Model* class used with functional API
- `Sequential`

```

- from keras.models import Sequential
- model = Sequential()
  * the keras.models.Sequential class is a wrapper for the neural network model
- the class Sequential() provides common functions like
  * hist=model.fit(X,y,nb_epoch=10,validation_split=0.2)
  * model.compile('adam', 'categorical_crossentropy', ['accuracy'])
  * model.compile(loss='mse',optimizer='adam')
  * model.fit(X,y,nb_epoch=1,batch_size=8,validation_data=(Xv,yv), shuffle=True)
    · h = fit() returns a history object, where h.history contains loss and val_loss as keys,
    · h.history['loss'] & h.history['val_loss'] can be plotted as a function of epochs
  * model.evaluate(x, y, batch_size=32, verbose=1, sample_weight=None)
  * model.add() adds a layer, see next
  * model.save(fn)

```

- Layers:

```

- from keras.layers.core import Dense, Activation, Flatten, Dropout
- from keras.layers import Input, Lambda
- to add input to model type
  model.add(Flatten(input_shape=(32, 32, 3)))
- to normalize input, type
  model.add(Lambda(lambda x: (x / 255.0) - 0.5))
  * more generally, Lambda layers can be used to create arbitrary functions that operate on each input, as it passes
  through the layer
- to add a second FC layer, type
  model.add(Dense(100))
- to add various activations, type
  model.add(Activation('relu'))
  model.add(Activation('softmax'))
- to add dropout, type
  model.add(Dropout(0.5))
- Keras infers the shape of subsequent layers after the first layer

```

- Convolution & pooling:

```

- from keras.layers.convolutional import Convolution2D
- from keras.layers.pooling import MaxPooling2D
- model.add(Convolution2D(32, 3, 3, input_shape=(32, 32, 3)))
  where 32 are the features, and 3, 3 is the kernel
- model.add(MaxPooling2D((2,2)))

```

- Datasets:

```

from keras.datasets import cifar10
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

```

- Model

```

- from keras.models import Model
- from keras.layers import Input, Dense
- a = Input(shape=724,)
- b = Dense(32)(a)
- model = Model(inputs=a, outputs=b)

```

- Callback:

```

- a callback is a set of functions to be applied at given stages of the training procedure
- can use callbacks to get a view on internal states and statistics of the model during training
- can pass a list of callbacks (as the keyword argument callbacks) to the .fit() method of the Sequential
  or Model classes
- the relevant methods of the callbacks will then be called at each stage of the training
- ModelCheckpoint() saves the model after every epoch
  keras.callbacks.ModelCheckpoint(filepath, monitor='val_loss',
    save_best_only=False, save_weights_only=False, period=1)
- EarlyStopping() stops training when a monitored quantity has stopped improving
  keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0, patience=0)

```